

Rui Pan

ruipan.xyz/ | github.com/ruipeterpan | ruipan@princeton.edu | [Google Scholar](#)

EDUCATION

Princeton University

Ph.D. in Computer Science

- M.A. in Computer Science, May 2025 (conferred en route to Ph.D.)
- Advisor: Prof. Ravi Netravali

Princeton, NJ, USA

Aug 2022 – May 2027/8 (Expected)

University of Wisconsin-Madison

B.S. in Computer Science and B.S. in Mathematics

- GPA: 3.96/4.00
- Advisor: Prof. Shivaram Venkataraman

Madison, WI, USA

Sep 2018 – Dec 2021

RESEARCH INTERESTS

I am broadly interested in the intersection between systems, networks, and machine learning. My recent work has focused on **efficient systems and algorithms for ML/LLM inference**, and has been published at top systems and ML venues including SOSP, NSDI, MLSys, NeurIPS, HotNets, and IEEE ToN.

FIRST-AUTHOR PUBLICATIONS (*EQUAL CONTRIBUTIONS)

[5] **Fail Fast, Win Big: Rethinking the Drafting Strategy in Speculative Decoding via Diffusion LLMs.** **Rui Pan**, Zhuofu Chen, Hongyi Liu, Arvind Krishnamurthy, Ravi Netravali. In *submission*, arXiv preprint available (2025).

[4] **SpecReason: Fast and Accurate Inference-Time Compute via Speculative Reasoning.** **Rui Pan**, Yinwei Dai, Zhihao Zhang, Gabriele Oliaro, Zhihao Jia, Ravi Netravali. In *The Thirty-Ninth Annual Conference on Neural Information Processing Systems (NeurIPS '25)*. **Spotlight Paper at The NeurIPS 2025 Workshop on Efficient Reasoning**.

[3] **Marconi: Prefix Caching for the Era of Hybrid LLMs.** **Rui Pan**, Zhuang Wang, Zhen Jia, Can Karakus, Luca Zancato, Tri Dao, Yida Wang, Ravi Netravali. In *The Eighth Annual Conference on Machine Learning and Systems (MLSys '25)*. **Outstanding Paper Award, Honorable Mention. Being integrated into SGLang!**

[2] **Apparate: Rethinking Early Exits to Tame Latency-Throughput Tensions in ML Serving.** Yinwei Dai*, **Rui Pan***, Anand Iyer, Kai Li, Ravi Netravali. In *The 30th Symposium on Operating Systems Principles (SOSP '24)*.

[1] **Efficient Flow Scheduling in Distributed Deep Learning Training with Echelon Formation.** **Rui Pan***, Yiming Lei*, Jialong Li, Zhiqiang Xie, Binhang Yuan, Yiting Xia. In *The 21st ACM Workshop on Hot Topics in Networks (HotNets '22)*.

ALL PUBLICATIONS

[11] **Optimizing Mixture-of-Experts Inference Latency Combining Model Deployment and Communication Scheduling.** Jialong Li, Shreyansh Tripathi, Lakshay Rastogi, Yiming Lei, **Rui Pan**, Yiting Xia. In *The IEEE/ACM Transactions on Networking (ToN '26)*.

[10] **METIS: Fast Quality-Aware RAG Systems with Configuration Adaptation.** Siddhant Ray, **Rui Pan**, Zhuohan Gu, Kuntai Du, Shaoting Feng, Ganesh Ananthanarayanan, Ravi Netravali, Junchen Jiang. In *The 31st Symposium on Operating Systems Principles (SOSP '25)*.

[9] **Mowgli: Passively-Learned Rate Control for Real-Time Video.** Neil Agarwal, **Rui Pan**, Francis Y. Yan, Ravi Netravali. In *The 22nd USENIX Symposium on Networked Systems Design and Implementation (NSDI '25)*.

[8] **Humanity's Last Exam.** Center for AI Safety, Scale AI. ArXiv preprint available (2025).

[7] **Improving DNN Inference Throughput Using Practical, Per-Input Compute Adaptation.** Anand Iyer, Mingyu Guan, Yinwei Dai, **Rui Pan**, Swapnil Ghandi, Ravi Netravali. In *The 30th Symposium on Operating Systems Principles (SOSP '24)*.

- [6] **Shockwave: Fair and Efficient Cluster Scheduling for Dynamic Adaptation in Machine Learning.**
Pengfei Zheng, Rui Pan, Tarannum Khan, Shivaram Venkataraman, Aditya Akella. In *The 20th USENIX Symposium on Networked Systems Design and Implementation (NSDI '23)*.

EXPERIENCE

Google, SystemsResearch@Google <i>Student Researcher</i>	Jun 2025 – Dec 2025 <i>Sunnyvale, CA, USA</i>
• Mentor: Prof. Arvind Krishnamurthy. • Topic: Systems for efficient LLM inference. • Work in progress: Efficient inference of diffusion-based large language models.	
AWS AI, Machine Learning Systems Team <i>Applied Scientist Intern</i>	May 2024 – Dec 2024 <i>Santa Clara, CA, USA</i>
• Mentors: Dr. Zhen Jia, Dr. Zhuang Wang, and Dr. Can Karakus. • Topic: Systems for efficient LLM inference. • Researched efficient prefix caching for hybrid LLMs (e.g., Qwen3-Next, Gemini) that combine Attention layers with subquadratic layers (e.g., Mamba, Sliding Window Attention). • Implemented and open-sourced Marconi, a radix tree-based prefix management system; contributed to its ongoing integration into SGLang (a popular LLM inference framework with 19k+ GitHub stars). • Improved efficiency by up to 34.4× compared to state-of-the-art prefix caching systems (e.g., vLLM and SGLang). • Received the Outstanding Paper Award (Honorable Mention) at MLSys 2025.	
Princeton University <i>Graduate Research Assistant</i>	Aug 2022 – Present <i>Princeton, NJ, USA</i>
• Advisor: Prof. Ravi Netravali. • Topic: Systems for efficient ML/LLM inference. • SpecReason: Low-latency reasoning model inference paradigm that generalizes speculative decoding by relaxing per-token equivalence into semantic equivalence, achieving complementary efficiency gains. • METIS: Low-latency RAG system that jointly adapts RAG configurations per query and schedules them. • Apparate and E3: Low-latency ML and LLM inference via early exiting. • Mowgli: Offline reinforcement learning for rate control adaptation in video conferencing. • Ongoing projects: Efficient multi-round agents, efficient multimodal model inference, efficient deep research agents.	
Max Planck Institute for Informatics, Network and Cloud Systems Group <i>Research Intern</i>	Feb 2022 – Aug 2022 <i>Saarbrücken, Germany</i>
• Advisor: Prof. Yiting Xia. • Topic: Networked systems for efficient ML and LLM training. • Researched network traffic patterns of parallelization strategies in distributed deep learning training. Proposed a network abstraction for flow scheduling in large-scale ML/LLM training clusters.	

PROFESSIONAL ACTIVITIES

Reviewer:

- Conferences: MLSys 2026
- Workshops: NeurIPS 2025 Workshop on Efficient Reasoning, KDD 2025 Workshop on Inference Optimization for Generative AI
- Journals: Transactions on Machine Learning Research, IEEE Transactions on Parallel and Distributed Systems, IEEE Transactions on Networking, IEEE Transactions on Cloud Computing, IEEE Transactions on Mobile Computing, IEEE Transactions on Computers, IEEE Transactions on Machine Learning in Communications and Networking, IEEE Communications Letters, Computer Networks, Future Generation Computer Systems

Artifact Evaluation Committee: MLSys '23, OSDI '23, ATC '23

Student Volunteer: CMMRS '22, N2Women@SIGCOMM '22

Teaching Assistant: COS 316 Principles of Computer System Design Fa'23, COS 598D Systems and Machine Learning Sp'24

INVITED TALKS

- Distributed Systems Lab, University of Pennsylvania, Dec 2024
- IEEE EDS/EPS Lunch Series: Frontier AI, UCLA, Mar 2025
- ML Systems Group, UCSD, May 2025
- KDD 2025 Workshop on Inference Optimization for Generative AI, Aug 2025

AWARDS

- MLSys 2025 Outstanding Paper Award, Honorable Mention (top 2.5% of all submitted papers)
- MLSys 2025 Travel Grant (\$1,000)
- NeurIPS 2025 Workshop on Efficient Reasoning, Spotlight Paper (top 18.5% of all accepted papers)

GRADUATE-LEVEL COURSEWORK

- COS 518: Advanced Computer Systems
- COS 568: Systems and Machine Learning
- COS 597K: Systems for Serving Generative AI
- COS 510: Programming Languages
- COS 529: Advanced Computer Vision
- ECE 539B: Security and Performance Challenges in Networked Systems
- COMP SCI 736 (taken during undergraduate at UW-Madison): Advanced Operating Systems
- COMP SCI 759 (taken during undergraduate at UW-Madison): High Performance Computing for Applications in Engineering
- COMP SCI 744 (audited during undergraduate at UW-Madison): Big Data Systems

Last Updated: Dec 2025